

Plates

There are just 40 plate images for training which is not much, so I considered the following approach. First I undertook some fine-tuning to try and improve my score.

1. The first thing i did was to use a pre-trained model InceptionResNet V2.
2. I froze the model preserving the learnings from the model and unfroze the model with the following

	froze	unfrozen
freeze	59274	3494

3. The finding was that it was better to freeze the model: `base_model.trainable = False`
4. The next thing was to update hyper parameters. The approach was to change one hyperparameter and begin training again.
5. The first hyper parameter was batch_size and i undertook the following with these results

batch_size	16	56586
batch_size	32	59274
batch_size	64	45564

6. The take away is that there is an optimal batch_size for a model that has an impact on the performance of the models prediction.
7. There is a relationship between the batch_size and learning_rate and its worth adjusting this value if you change a batch_size. I found the following

learning_rate	0.0001	59274
learning_rate	0.005	54435

8. The result showed that learning rate has an impact on the models performance when generalizing on test data. Again this is a case of adjusting one parameter and retraining the model and observing the results.

9. Next up was epochs, the number of times we go through each cycle key for training.

epochs	50	59274
epochs	100	54564

10. A key observation is that this actually becomes irrelevant if you use early stopping. This stops training before the model begins to overfit. In most occasions the model will call early stopping before it completes all epochs.
11. For small datasets epochs from 50-100 are recommended for better learning as the model sees the same data repeatedly. For large data sets it is recommended to use 10 to 30.
12. Without early stopping if you have too many the model can overfit.
13. The next step was to include a `validation_generator` which is best practice, something that I did not include. It's worth noting without validation you miss out on some metrics such as `val_accuracy` and `val_loss` and cannot use early stopping.
14. A key parameter in early stopping is patience. I adjusted this parameter with the following findings

patience	5	52419
patience	10	58870
patience	15	56989

15. So patience is just giving training a chance e.g. if there are 5 epochs and no improvement in say `val_loss` and say it starts getting worse training is terminated early, and the weights are restored from the best epoch.
16. This stops overfitting. Using `verbose=1` tells you when it terminates and why.
17. The last part of hyper tuning was using `steps_per_epochs`. This is in general the number of batches an epoch will see before completing and moving to the next epoch. Again I found this influenced the final score when using predict with the model.

steps_per_epochs	100	59946
steps_per_epochs	50	42607
steps_per_epochs	150	32215

18. Again with making an adjustment, then retraining, steps per epoch had some influence over the models prediction.
19. If too small the model cannot see enough data per epoch and results in poor generalization.
20. if too large then the model overfits and gives poor predictions.
21. I took it as far as i could. Then i considered the pre-trained model.
22. It turns out that the inceptionResNetv2 model is super complex, great for large images with complex data sets and tends to overfit with small data sets.
23. My dataset is 40 images so obviously changing to a less complex, small data friendly model makes sense.
24. I employed Resnet50 a simple model that works great with small data sets so i used this.

Before

```
base_model = applications.InceptionResNetV2(weights='imagenet',
include_top=False,
#
input_shape=(image_size,image_size, 3))
```

After

```
base_model = applications.ResNet50(weights='imagenet', include_top=False,
input_shape=(image_size,image_size,3))
###
```

25. I decide to keep the pre-models weights from the convolutional layers to dense (the top) frozen to preserve the learnings from the new model.
26. There was one more adjust required involving the flattening layer in the top layer of the CNN.
27. I used

```
x = Flatten()(x)
```

28. This is a complex function that flattens the output to pass this on to the dense layer. However this layer is great for complex models with substantial data, a poor fit for a small model with 40 images. This is the axes between the convolutional layers and the connected layer. I found something better fitting with a small scale model such as mine and used this.

```
x = GlobalAveragePooling2D()(x)
```

29. This compress the feature maps to small vectors and is optimal for small datasets and reduces overfitting. This passes on to the dense layer (connected layer)
30. The results

ResNet 50 + GlobalAveragePooling	63440
inceptionResnet + Flatten	57930

Steps_per_epoch

31. Ok so i changed to pre-trained ResNet 50 and decided to perform some tuning of the model.
32. I read that there is a formula e.g. images / batch size = steps per epoch.
33. With the existing batch size following best practice e.g. 16, 32, 64 using 32 for size of data set i found the following
34. The changed resulted in a marginal improvement.

steps_per_epoch	100	63440
steps_per_epoch	7	64784

Note:

I built a learning rate scheduler that is used to create a **learning rate scheduler** that reduces the learning rate when the performance of the model plateaus, which helps the model converge more effectively.

```
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=5,
min_lr=1e-6)
```

I therefore do not need to manually adjust the learning rate and number of epochs is academic with early stopping with patience set a 10 for best practice.

35. The next option would be around unfreezing the pre-trained model. This would result in the pre-trained weights being changed during training.
36. I could unfreeze a subset of the pre-trained base convolutional layers, to try not to go to far.
37. I used the following code to unfreeze just 10 layers before the output (this is part of the pre-trained layers). Why are we doing this?

```
for layer in base_model.layers[-10:]:
    layer.trainable = True
```

frozen	64784	
unfreeze 10 layers	63306	

38. The performance decreased!

39. **Small Datasets:** With a small dataset, the pre-trained weights might already be well-suited for your task, and unfreezing layers could lead to overfitting. In such cases, the model might benefit more from keeping the pre-trained layers frozen.

40. The result was that the model started to overfit slightly and the performance decrease slightly.

41. This highlights the pre-trained weights were already well suited for the task and allowing the 10 top layers before output to be adjusted resulted in overfitting.